

Vector databases index high-dimensional embeddings so that nearest-neighbor similarity search stays